CiteWise Sample Document

CiteWise is a lightweight PDF ingestion system.

It extracts text, detects scanned pages, and uses OCR.

The pipeline preserves page numbers and document metadata.

Why Chunking Matters

Chunking splits long documents into smaller passages.

Each chunk retains its source page for future citations.

Overlap can preserve context between adjacent chunks.

Project Test Facts

Project: CiteWise

Stage: PDF ingestion

Text extraction: PyMuPDF

OCR engine: Tesseract

Default chunk size: 1000 characters

Default overlap: 150 characters